

R₅₀ Estimation and Comparison: Final Report

Choosing a Model for Estimating R₅₀

When first looking at this problem and the data I was given I knew that a logistic regression could be used to estimate R₅₀. I decided that a good estimate of R₅₀ is the range where there is a 50% probability a target will be detected since R₅₀ is the range where 50% of the targets will be detected. However, a prediction using a logistic model can't estimate R₅₀ directly. Logistics output a probability given a continuous variable, whereas what I needed was the opposite. I looked to see if there was a way to create a model that would be able to do this "out-of-the-box" but I found nothing and decided to stick with the logistic. I chose to create a logistic regression using the R package **parsnip** since prediction output is in a dataframe and I could obtain confidence intervals easily from it as well.

Solving for R₅₀ and Confidence Intervals

Since my model couldn't estimate R₅₀ directly, I created a function `solve_R50` that could use the model to solve for R₅₀ and the confidence intervals under each condition.

In the function below the argument `R50_mod` is the **parsnip** logistic model and `turbines` is a binary vector indicating the condition to predict under. The argument `type` is one of "pred", "lower" or "upper". "pred" gives the predicted R₅₀ value. "lower" or "upper" give the corresponding side of the confidence interval.

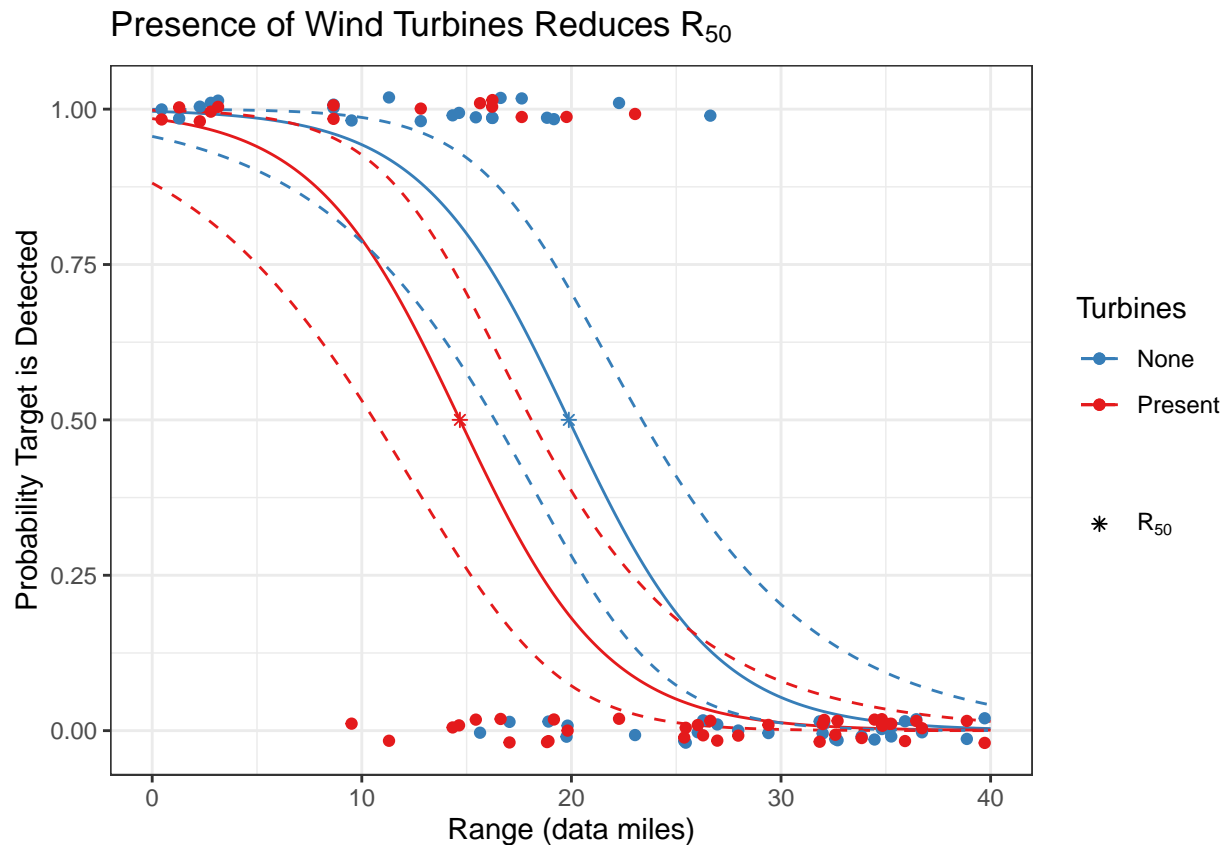
```
solve_R50 <- function(R50_mod, turbines, type) {  
  
  # ... skipped logic for `type` argument  
  
  # create a function that we will find the root of using `uniroot`  
  R50_root <- function(range, turbines) {  
    predict.model_fit(R50_mod,  
                      tibble(range = range, turbines = turbines),  
                      # pred_type variable determines if we are solving  
                      # for R50 or a confidence interval  
                      type = pred_type  
    ) %>%  
    # pick the column we need from predict.model_fit  
    # subtracting .5 creates a root where we want it  
    pull(pred_col) - .5  
  }  
  
  # solve for each `turbines` condition  
  map_dbl(turbines,  
    ~ uniroot(R50_root,  
              interval = range(R50_mod$fit$data$range),  
              turbines = .x)$root)  
}
```

This is how I used the solver to obtain R_{50} and the confidence intervals for each turbine condition. Below is the information presented in table form and graphically.

```
R50_dat <-
  tibble(turbines = 0:1) %>%
  mutate(R50 = solve_R50(R50_mod, turbines, "pred"),
         lower = solve_R50(R50_mod, turbines, "lower"),
         upper = solve_R50(R50_mod, turbines, "upper"))
```

R_{50} Estimation and Confidence Intervals

| Turbines | R_{50} | Lower | Upper |
|----------|----------|-------|-------|
| None | 19.88 | 16.35 | 23.37 |
| Present | 14.68 | 10.59 | 18.09 |



Identify the Impact of the Presence of Wind Turbines on R_{50}

We can determine whether the presence of wind turbines affects R_{50} by looking at the impact the `turbines` term has on the model. Below are the estimates of the model's terms. Since the `turbines` term is significant and it represents a horizontal shift in the two curves as shown above, the estimated R_{50} will be significantly different under the two conditions.

R_{50} Model Terms

| term | estimate | std.error | statistic | p.value |
|-------------|----------|-----------|-----------|----------|
| (Intercept) | 5.64 | 1.29 | 4.38 | 0.000012 |
| range | -0.28 | 0.06 | -4.70 | 0.000003 |
| turbines | -1.47 | 0.69 | -2.12 | 0.033855 |

Estimation of power

In order to calculate power I needed to estimate the standard deviation of R_{50} . Unfortunately, the confidence intervals were not symmetric after translating them to be in terms of range. To be conservative in the power calculations I choose the largest margin to determine the standard deviation.

 R_{50} Confidence Margins

| Turbines | R_{50} | Lower | Upper | Lower Margin | Upper Margin |
|----------|----------|-------|-------|--------------|--------------|
| None | 19.88 | 16.35 | 23.37 | -3.53 | 3.50 |
| Present | 14.68 | 10.59 | 18.09 | -4.09 | 3.41 |

If we assume there are two standard deviations from R_{50} to the interval bounds, then the largest the standard deviation could be is about 2. Since the difference in R_{50} is about 5 data miles less under conditions when there are wind turbines verses no turbines, then the R_{50} s are about 2.5 standard deviation from each other. In other words, the effect size is about 2.5 data miles. Using this information, I created the plot below.

Power of finding differences in R_{50} under different turbine conditions

Effect size in this sample was ~ 2.5

